

Chapter 1

Silhouette Execution

In the previous two chapters, we identified how and why multiple executions of the same Linux service can diverge in behavior. In light of our results, this chapter introduces a novel design strategy that aims to mitigate the bootstorm problem: *silhouette* execution (Section 1.1). For the sake of evaluation, this chapter presents some design sketches for silhouette execution for user-mode programs (Section 1.2). It also describes the simple simulation techniques we used to model the effectiveness of silhouette execution using our design sketches (Section 1.3). From the results of our simulations, we present strategies to improve the effectiveness of silhouette execution (Section 1.4) and evaluate them (Section 1.5). Our simulations show that the proposed designs can be successful in mitigating the bootstorm problem.

1.1 What is *Silhouette* execution?

The previous few chapters showed that distinct executions of the same Linux service generally execute the same number of instructions across identical VMs when they start up. If the number of booting VMs for the same physical host is large (as is the case in VDI deployments), then executing many instructions over and over again during boot represents a waste of scarce CPU cycles. *Silhouette* execution, in essence, targets redundant execution of the same instructions across distinct but identical virtual machines in order to reduce CPU pressure in bootstorm scenarios.

As shown in Figure 1-1, silhouette execution is analogous to page sharing: both design ideas aim to use hardware resources effectively to improve VM density per host in virtualization scenarios. While page sharing reduces pressure on the memory subsystem by identifying and refactoring overlapping memory contents, silhouette execution identifies overlapping instruction streams and refactors execution to reduce pressure on the CPU. Like memory overcommit, the ultimate aim is to allow a host to support VMs that together require more hardware resources than are really available in the host.

To the best of our knowledge, silhouette execution is a novel design idea that has not been suggested or implemented before. To study whether this approach can be effective in reducing CPU pressure in concurrently booting VMs, we present some design sketches for implementing silhouette execution for Linux services in the rest of this chapter. Admittedly, user-mode instruction streams from Linux services capture a subset of instructions executed by a booting VM. However, we focus on Linux services as a first step in studying the feasibility of silhouette execution. After all, as outlined Section ??, booting VMs can saturate host CPUs when they launch many identical user-space processes. For a complete solution, proposed design sketches need to be generalized to the execution of entire VMs themselves; this would require us to precisely identify the execution differences that can arise from all software layers inside a VM.

1.2 Silhouette Execution for Linux Services

For user-mode Linux services, our proposed design skeletons for silhouette execution use information recorded from one executing program – the *leader* – to bypass execution steps of subsequent instances of the program – the *silhouettes*. Ideally, the leader executes all instructions from the program, while the silhouettes execute a much smaller subset of these instructions.

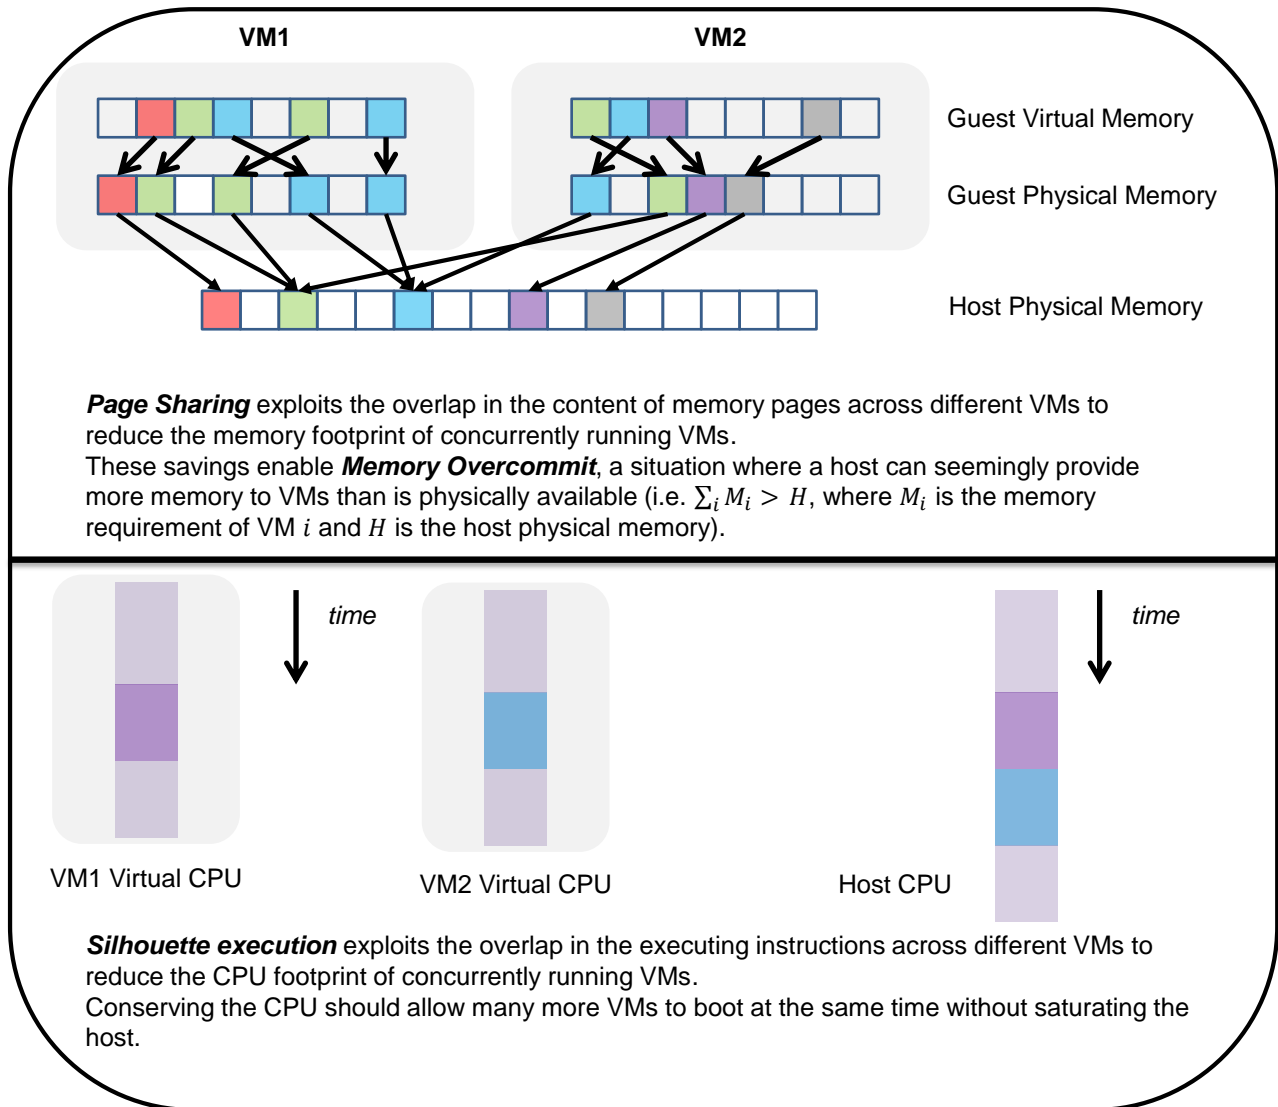


Figure 1-1: Silhouette execution is analogous to Page Sharing.

To maintain correctness in user-space, silhouettes need only execute:

- instructions that can potentially cause the leader’s execution to differ from the silhouettes (e.g. the `rdtsc` instruction);
- instructions that propagate any earlier differences in execution;
- instructions that write to memory;
- system calls that have side-effects to entities outside a user-space program (e.g. the `read` system call mutates hidden operating system state associated with file descriptors).

If there are no differences between the leader’s execution and a silhouette’s, then the silhouette would only execute the system calls and store instructions executed by the leader until the login screen is shown. Executing all the memory writes from the leader in the silhouettes ensures that the address space and memory contents of the two instances evolves in an identical manner. Executing all the system calls with the same arguments also ensures that the silhouette’s execution is *faithful* to its semantics, that is, the side-effects to the underlying operating system are maintained till the end. After we restore the register contents, a silhouette can simply continue execution independently of the leader. When the number of system calls and memory writes is a small fraction of the leader’s execution, this approach can theoretically reduce the stress placed on the host CPU.

More generally, when there are instructions with conflicting side-effects in the leader and a silhouette, then these instructions need to be executed by each silhouette independently. This ensures that the silhouette’s execution semantics are retained. It is not known *a priori* which instructions in the two instances of the same program will behave differently or not. Our detailed analysis of the various interfaces between application programs and the operating system allows us to identify such potential sources of execution divergence via dynamic program inspection.

Note that silhouette execution for user-space programs is fundamentally different from record-and-replay approaches because it does not semantically alter subsequent executions of a program by emulating the leader’s execution. In fact, silhouettes are independent executions of a program that can potentially branch from the leader’s execution at any point.

The next few subsections outline a few design sketches for implementing silhouette execution on individual user-space programs such as Linux services. We have not implemented these designs. Instead, we present them here to evaluate the effectiveness of silhouette execution in user-space.

1.2.1 *Precise Silhouetting*

Here is a simple design that uses silhouette execution to refactor execution in a user-mode program:

1. We run one program instance – the *leader* – slightly ahead of all other program instances – the *silhouettes*.
2. Using dynamic inspection techniques on the leader, we
 - precisely identify instructions where other instances of the program could *potentially* diverge. We call these instructions *fork-points*.
 - collect an *execution signature* that summarizes the leader’s execution between successive fork-points. For a user-space program, this includes a record of memory operations and deterministic system calls.
3. When a leader reaches a fork-point, it sends its own execution signature from the previous fork-point (or the beginning of the program) till the current fork-point to all other silhouettes.
4. The silhouettes do not execute all the instructions that the leader executes. In fact, each silhouette bypasses execution between two fork-points by executing only the memory operations and system calls from the execution signature sent by the leader, and restoring the register state at the end.

5. When a silhouette reaches a fork-point, it independently executes the *forking* instruction and observes its side-effects. The forking instruction may or may not behave differently in a silhouette than the leader.

- If the forking instruction does have different side-effects in a silhouette, the silhouette branches execution and executes instructions independently from that point onwards. We call this instruction an *active* fork-point.
- Otherwise, we call this instruction a *latent* fork-point. We return to step 4: the silhouette waits for the leader’s next execution signature for bypassing execution to the next fork-point.

We name this design *precise silhouetting* because it cannot tolerate any differences in execution between the multiple instances of a program: silhouettes completely branch execution after executing a forking instruction that disagrees with the leader (i.e. an active fork-point). Our description of precise silhouetting implies that the leader executes concurrently with the silhouettes – albeit slightly ahead – but this is not necessary. This approach would work even if we run the leader to completion before we run any silhouettes. The silhouettes, of course, would only execute system calls and memory operations between successive fork-points that the leader recorded earlier until execution diverges.

1.2.2 *Optimistic Silhouetting (excluding control flow)*

Optimistic silhouetting essentially follows the same overall design principles as precise silhouetting, except that it allows silhouettes to tolerate minor execution differences before branching execution completely. In this design:

1. The leader executes slightly ahead of the silhouettes. The leader identifies fork-points and sends execution signatures to silhouettes. The silhouettes bypass execution by only executing the load/store instructions and system calls made by the leader, and restoring register contents at the end of each fork-point. This is what happens in precise silhouetting before the first active fork-point.

2. Unlike the previous design, when the leader reaches any fork-point, it always waits for the silhouettes to catch up with it. All the instances execute a forking instruction in sync and compare its side-effects.
3. If a forking instruction has different side-effects in a silhouette than the leader (i.e. at an active fork-point):
 - the silhouette does not immediately branch execution completely;
 - the leader tracks the register or memory values that are written differently in the multiple instances by marking them as *tainted*;
 - the leader treats any subsequent instructions that read tainted values as fork-points as well;
 - the silhouettes do not overwrite the values in any tainted registers with those contained in the leader’s execution signature.
4. When fork-points become too frequent, or when control flow diverges (e.g. a tainted value is compared to a constant to determine control flow), a silhouette starts executing instructions independently and branches off from the leader.

This approach does require that that the leader and its silhouettes execute fork-points at the same time and communicate their results. This is necessary so that the leader can identify subsequent instructions that propagate any earlier nondeterminism (e.g. read a tainted value) as fork points.

1.2.3 *Optimistic Silhouetting (including control flow)*

This design is similar in essence to the version of *optimistic silhouetting* described above, but it can also tolerate minor control flow differences between the leader and the silhouettes. In this design:

1. As before, the leader and the silhouettes must execute fork-points concurrently. The leader transmits execution signatures to silhouettes, and uses dynamic taint propagation to tolerate minor differences in instruction side-effects.

2. Unlike before, silhouettes do not branch off permanently from the leader at the sign of the first control flow divergence:
 - The leader uses dynamic instrumentation to create a dynamic *control flow graph* for the program execution.
 - When the leader and a silhouette reach a control flow fork-point with divergence (e.g. when a tainted value is read to determine control flow), the leader uses the dynamic control flow graph to determine the *immediate post-dominator* of the block where execution has diverged.
 - The silhouette branches off temporarily (rather than permanently) and executes independently to the point of control flow convergence. The silhouette and the leader log any memory values written or any system calls made during this *branch interval*.
 - The leader and the silhouette compare their current register state, along with the system calls made or memory values written during the branch interval.
 - An *analysis engine* figures out whether the two executions are reconcilable or not based on what happened in the branch interval.
3. If the two executions can be reconciled, any conflicting state (e.g memory addresses or register values are marked by the leader) as tainted, and the silhouettes start waiting for execution signatures from the leader again.
4. If the two executions cannot be reconciled, or when fork-points become too frequent, execution branches permanently.

Reconciling Executions

The notion of whether two distinct execution strands from a branch interval can be reconciled is a new one. If two instances do not execute any system calls or memory operations during the branch interval, then execution can be simply reconciled by marking any different register values as tainted. If two instances do execute some

memory load/store operations, then different memory contents can be marked as tainted as well to reconcile them.

If the two instances make different system calls during the branch interval, then execution may or may not be reconcilable. If the system calls are stateless (e.g. `time`), then execution can clearly be reconciled. On the other hand, if one execution strand makes system calls that change the operating system state, then the leader must know how to identify any subsequent system calls that depend on the changed state. For instance, if a silhouette does an extra `open` to a file in the branch interval, the leader must treat each subsequent `open` as a fork point, because the returned file descriptors will now be different. The leader may have previously assumed that all files to be opened were present on the identical VMs and thus not treated the `open` system call as a fork-point by default.

There is a clear trade-off in the complexity in the dynamic instrumentation layer in the leader that tracks dependencies across system calls and the extent to which we can prolong silhouette execution. For simplicity, we will assume that if any instance executes a system call in its branch interval that mutates some external state, then the executions are irreconcilable.

1.3 Evaluation Scheme

The data collection scheme described in Chapter ?? does not actually implement silhouette execution in user-space because multiple instances execute all application instructions independently. This section describes how we can still mathematically simulate silhouette execution by comparing execution traces from our data collection scheme.

There are several factors to consider in determining the effectiveness of silhouette execution. Ideally, we would like the following conditions to be true:

- The first forking instruction with conflicting side-effects (i.e. the first active fork-point) should occur as late as possible into the leader's execution. This is especially important for precise silhouetting, because silhouettes branch-off permanently after the first active fork-point.
- The number of fork-points should be much smaller than the total number of instructions executed by the leader. All the instances have to analyze the side effects of each forking instruction to determine whether execution has diverged or not, which represents a serious design overhead.
- The number of active fork-points must be small. Fewer active fork-points would create fewer tainted memory locations, and thus reduce the overhead in the leader associated with dynamic taint propagation.
- The number of control flow divergences should be very small. Any control flow divergences should preferably be short-lived, and have few system calls or memory operations. This reduces the overhead associated with reconciling executions after branch intervals and creating dynamic control flow graphs in the leader.
- The fork-points must be separated by very many instructions so that memory access logs can be compressed. We could forget intermediate values of memory locations and only remember their final values instead.
- Programs should have a high ratio of user-mode instructions to system-calls and memory operations so that silhouettes execute few instructions compared to the leader when they are bypassing execution.

1.3.1 Computed Metrics

For our data collection scheme, we can identify fork-points by simply parsing the traces collected by our Pin tool and looking for the sources of potential execution differences cataloged in the previous chapter. Once we can identify individual fork-points, we can compute:

- The number and distribution of fork-points – both latent and active – in a program,
- The number and distribution of control flow divergences in a program,
- The proportion of memory and system-call instructions between successive fork-points,
- Size estimates for execution trace files that need to be communicated between the leader and its silhouettes.

Using simple mathematical models, we can compute the number of user-space instructions the host CPU has to execute *without* silhouette execution (T_O), and the number of user-space instructions the host CPU has to execute under silhouette execution (T_S). We measure the advantage conferred by silhouette execution, A , as:

$$A(\vec{K}, N) = \frac{T_O}{T_S}. \quad (1.1)$$

A is parameterized by $\vec{K} = (k1, k2, k3...)$ and N . \vec{K} represents the constants associated with the overhead of various aspects of silhouette execution and N is the number of concurrently running instances of a program.

A value of $A > 1$ implies that silhouette execution is effective in reducing CPU overhead from concurrent program execution on the host in user-space. Generally, A should increase as N increases (holding everything else constant), and A should decrease as individual entries in \vec{K} increase (holding everything else constant). T_O is easily computed: $T_O = NI$, where N is the total instances of a program to be run, and I is the number of instructions each instance must execute. Typically, I is the

number of instructions necessary to model the start of a program or a service. For many Linux programs, a few iterations of the main scheduler loop of the program is sufficiently representative of execution before a login screen is shown. The value of T_S depends on which version of silhouette execution is being used.

Precise Silhouetting

Given multiple traces, instructions that are in the common prefix (P) broadly represent savings from precise silhouetting. Figure 1-2 summarizes how T_S and \vec{K} can be computed for precise silhouetting.

Optimistic Silhouetting (Excluding Control Flow)

Instructions in the longest common subsequence (LCS) of multiple traces *before* a control flow divergence broadly represent the savings from this variant of optimistic silhouetting. Figure 1-3 summarizes how T_S and \vec{K} can be computed for this variant of optimistic silhouetting.

Optimistic Silhouetting (Including Control Flow)

Instructions in the longest common subsequence (LCS) of multiple traces before execution diverges permanently represent the savings from this variant of optimistic silhouetting. Figure 1-4 summarizes how T_S and \vec{K} can be computed for this design.

1.3.2 Caveats

Before we present our results, we note a few limitations of our methods for evaluating silhouette execution:

- We estimate the advantage (A) of silhouette execution on user-space programs purely in terms of the number of instructions executed on the host CPU. We do not model *latency* for silhouette execution. It would be interesting to study whether the delays introduced by dynamic inspection of program execution and inter-instance communication can eclipse the potential latency reduction

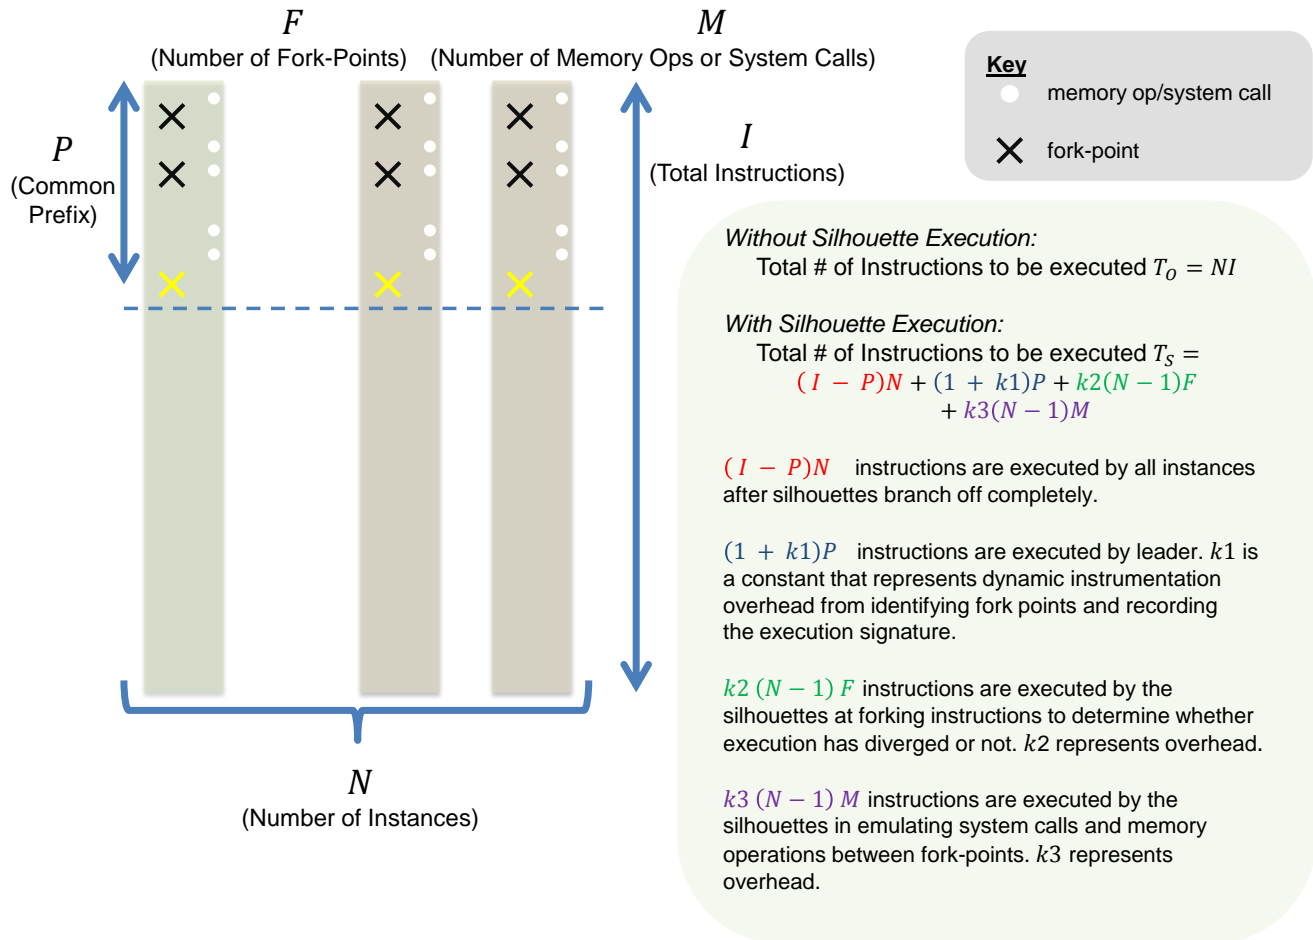


Figure 1-2: A simple way to model CPU overhead from precise silhouetting. We can compare the user-space instructions executed from running a program multiple times in the status quo versus the number of instructions executed when precise silhouetting is used for the same scenario. k_1 , k_2 and k_3 are constants that represent overheads associated with this approach.

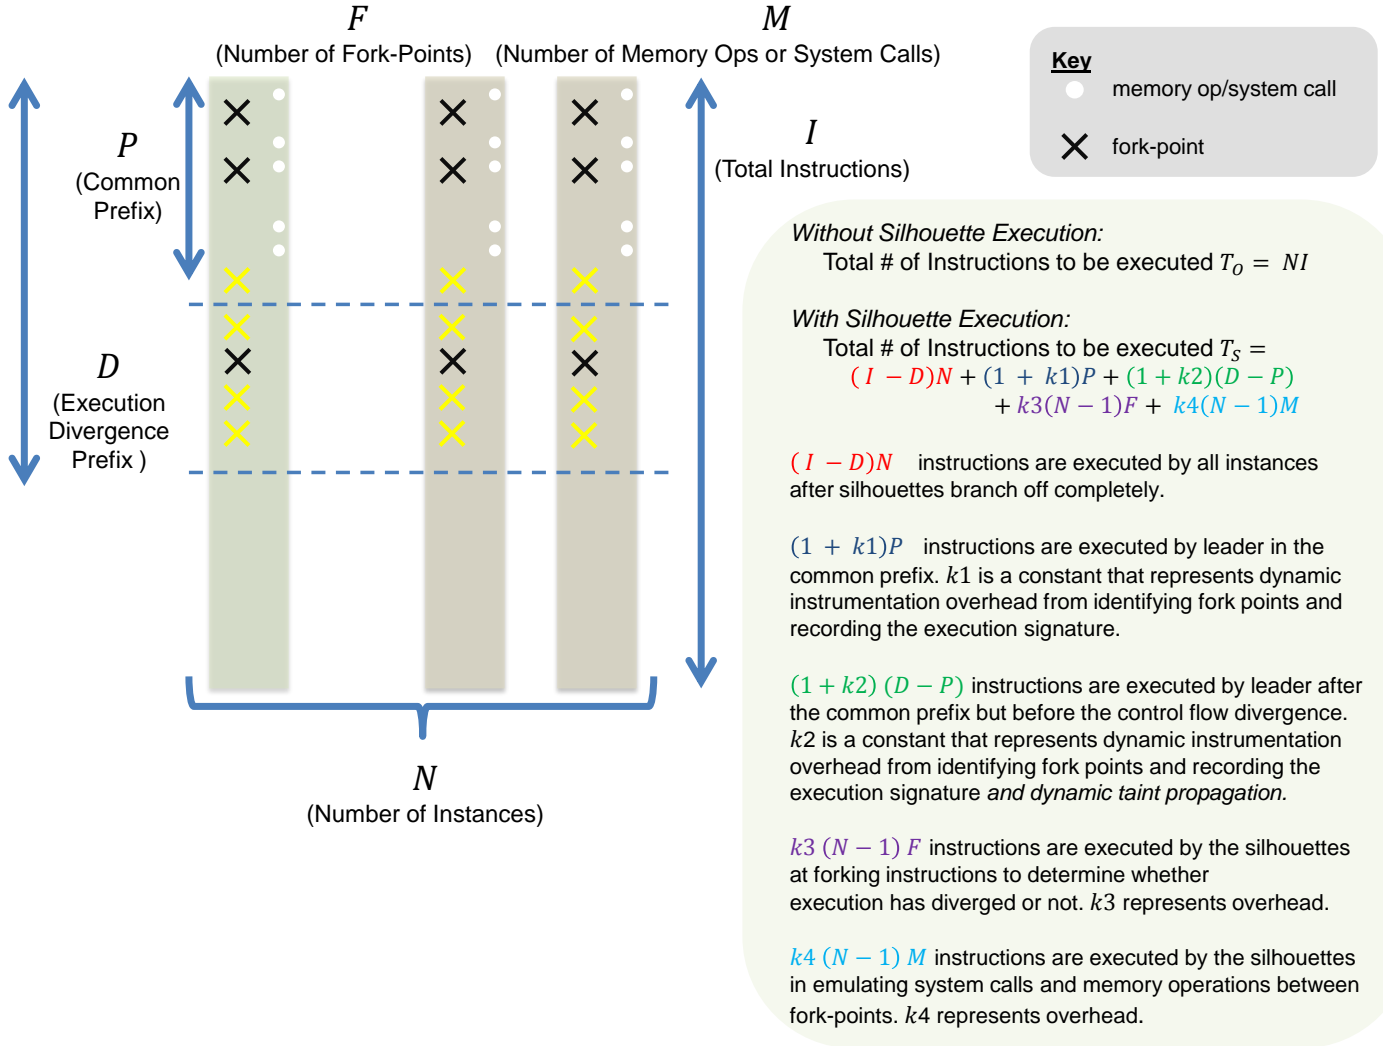


Figure 1-3: A simple way to model CPU overhead from optimistic silhouetting (excluding control flow). We can compare the user-space instructions executed from running a program multiple times in the status quo versus the number of instructions executed when optimistic silhouetting is used for the same scenario. k_1 , k_2 , k_3 , and k_4 are constants that represent overheads associated with this approach.

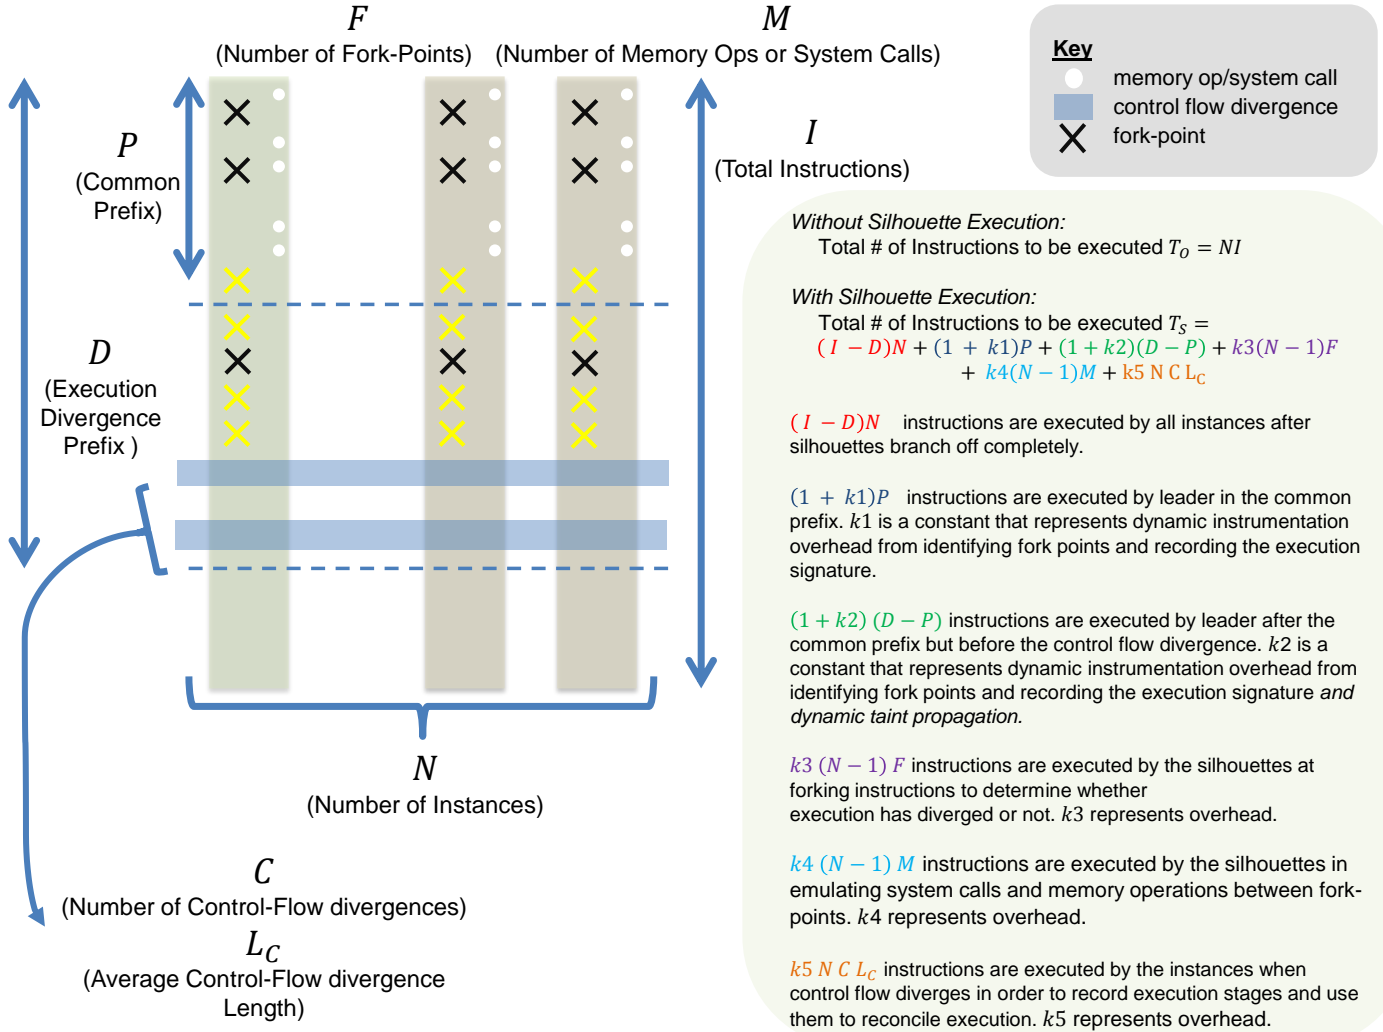


Figure 1-4: A simple way to model CPU overhead from optimistic silhouetting (including control flow). We can compare the user-space instructions executed from running a program multiple times in the status quo versus the number of instructions executed when optimistic silhouetting is used for the same scenario. k_1 , k_2 , k_3 , k_4 and k_5 are constants that represent overheads associated with this approach.

from reduced CPU load and bypassing execution in silhouette execution or not. In practice, the hypervisor layer rather than a dynamic instrumentation layer would implement silhouette execution, to reduce performance overhead.

- Our dynamic instrumentation tool can only inspect user-mode instructions of the main process hosting an application, so we cannot consider code executed by lower software layers or other children processes in computing A . Overall, the number of instructions we consider may be a fraction of all the instructions computed on the host CPU, which would add a dampening factor to our computed value of A .
- The values we use for \vec{K} are conservatively guessed, and we assume the overheads from various aspects of silhouette execution are linear in nature. These assumptions, while reasonable, may understate the CPU-load reduction practically attainable by silhouette execution in user-space.
- We do not factor the storage and I/O overhead associated with the transmission of execution signatures, though our experience suggest that signature files are typically very small (i.e. only a few megabytes) so they should fit in host memory.
- While we collect traces from many different silhouettes, we simply pick the worst trace (i.e. the one with the most difference from a leader) for computing T_S . Thus, our models may be overly conservative because they assume that *all* silhouettes are as different from a leader as the *worst* silhouette. This assumption simplifies the design and evaluation complexity related from the leader having to handle silhouettes with varying levels of divergence from the leader.

Despite these limitations, we believe that our model offers valuable insight into the feasibility of silhouette execution in user space because we instrument and evaluate silhouette execution on large user-mode instruction streams from Linux services, and conservatively factor in the possible overhead from silhouette execution.

1.3.3 Initial Results

Precise Silhouetting

Table 1.1 shows the results of modeling precise silhouetting on a few Linux services. For simplicity, we include all system calls in our definition of fork-points. This assumption increases the number of fork-points (F); it also increases the overhead associated with determining if execution has diverged or not after a fork-point because the inspection layer has to presumably understand the logic of each system call to determine its side effects ($k2$). However, this assumption reduces $k1$ because an overwhelming majority of instructions that only use register operands are easily excluded from fork-points. Instructions that are system calls are also easily identifiable; instructions with memory operands need their memory addresses to be compared to tainted addresses (e.g. `gs:0x14` or `gs:0x18`) to determine whether they are fork points or not. For our evaluation, we chose $k1 = 20$, $k2 = 1000$ and $k3 = 20$ as reasonably conservative values for the overhead constants.

Table 1.1: Preliminary Results from Modeling Precise Silhouetting

A , the advantage ratio is calculated by $\frac{T_O}{T_S}$. T_O is the total instructions computed in the status-quo whereas T_S is the total instructions computed under precise silhouetting. \vec{K} represents overhead constants; M is the number of system calls and memory operations made by the leader before the first active fork-point; F is the number of latent fork-points before the first active fork-point. $p = P/I$ the prefix ratio of the execution.

Program	N	I	T_O	p	\vec{K}	M	F	T_S	A
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3)	499	299	90000	1.11

Because p is very small on average for our sample of Linux services, this version of optimistic silhouetting does not offer any savings. The value of A is 0.40 on average, again representing a *degradation* on CPU load.

Optimistic Silhouetting (Excluding Control Flow)

Table 1.2 shows the results of modeling optimistic silhouetting (without control flow) on our sample of Linux services. As before, we include all system calls in our definition of fork-points. We use the same values of $k1 = 20$, $k3 = 1000$, $k4 = 20$ to model overheads from detection of fork-points in the leader, execution of forking instructions to distinguish active and latent fork-points in the silhouettes and bypassing execution in silhouettes respectively. We use $k2 = 2k1 = 40$ to model the additional overhead from taint propagation after the first divergence.

Table 1.2: Preliminary Results from Modeling Optimistic Silhouetting (Excluding Control Flow).

A , the advantage ratio is calculated by $\frac{T_O}{T_S}$. T_O is the total instructions computed in the status-quo whereas T_S is the total instructions computed under this variant of optimistic silhouetting. \vec{K} represents overhead constants; M is the number of system calls and memory operations made by the leader before the first active fork-point; F is the number of latent fork-points before the first active fork-point. $d = D/I$ the portion of the execution before the first control-flow divergence.

Program	N	I	T_O	d	\vec{K}	M	F	T_S	A
cupsd	10	10000	100000	0.04%	(1,2,3,4)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4)	499	299	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4)	499	299	90000	1.11

While d is not as small as p for these Linux services, this version of optimistic silhouetting does not offer any savings. The value of A is 0.50 on average, again representing a *degradation* on CPU load. This is largely because of a large number of fork points and the high overhead in tracking their side-effects.

Optimistic Silhouetting (Including Control Flow)

Table 1.3 shows the results of modeling optimistic silhouetting (including control flow) on our sample of Linux services. As before, we include all system calls in our definition of fork-points. We use the same values of $k1 = 20$, $k2 = 40$, $k3 = 1000$, $k4 = 20$

to model overheads from detection of fork-points in the leader, taint propagation, execution of forking instructions to distinguish active and latent fork-points in the silhouettes and bypassing execution in silhouettes respectively. We use $k5 = 20$ to model the additional overhead from execution reconciliation after a branch interval.

Table 1.3: Preliminary Results from Modeling Optimistic Silhouetting (Including Control Flow).

A , the advantage ratio is calculated by $\frac{T_O}{T_S}$. T_O is the total instructions computed in the status-quo whereas T_S is the total instructions computed under this variant of precise silhouetting. \vec{K} represents overhead constants; M is the number of system calls and memory operations made by the leader before the first active fork-point; F is the number of latent fork-points before the first active fork-point; C and L_C represent the number of control-flow divergences and their average length respectively. $d = D/I$ the portion of the execution before permanent execution divergence.

Program	N	I	T_O	d	\vec{K}	M	F	C	L_C	T_S	A
cupsd	10	10000	100000	0.04%	(1,2,3,4,5)	499	299	4	100	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4,5)	499	299	4	100	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4,5)	499	299	4	100	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4,5)	499	299	4	100	90000	1.11
cupsd	10	10000	100000	0.04%	(1,2,3,4,5)	499	299	4	100	90000	1.11

While d is larger when we allow for control flow divergences in these Linux services, this version of optimistic silhouetting does not offer any savings. The value of A is 0.30 on average, again representing a *degradation* on CPU load. This is again because of a large number of fork points and the high overhead in tracking their side-effects.

1.4 Improving Silhouette Execution

When we analyzed the execution traces collected by our data collection scheme, we found that several causes of execution divergence across our sample of Linux services were synthetic i.e. they were an artefact of external sources of nondeterminism rather than semantic differences in program execution. We thus decided to introduce a deterministic execution layer to our previous designs to eliminate as many sources of fork-points as possible, to improve the feasibility of silhouette execution.

1.4.1 Modified Data Collection Scheme

We modified our data collection scheme from Chapter ?? (shown in Figure ??) to simulate and evaluate silhouette execution in the bootstorm scenario.

As shown by Figure 1-5, we run one instance of the program – the leader – before all others. For the leader, we generate an execution log, as before, but we also augment the log by summarizing information about the sources of nondeterminism described in Section ?. For instance, we record information about signal timing, process IDs, time-related system calls in the trace signature file. Our Pin tool uses the trace signature of the leader to modify the instruction sequences executed by subsequent executions (the silhouettes) and reduce the number of fork-points as much as possible.

We run the leader to completion before the silhouettes. As before, we also do not bypass instructions in silhouettes so our modified data collection scheme still requires us to analyze these traces to simulate and evaluate silhouette execution.

1.4.2 Reducing Execution Differences across Instances

We now describe how dynamic instrumentation can be used to reduce the source of execution differences from the sources described in ?. While we modify the instances that execute after the leader, in practice many of our techniques eliminate fork-points altogether i.e. the leader can continue execution past the forking instruction, or avoid control flow differences by navigating around variability of I/O timing and latency.

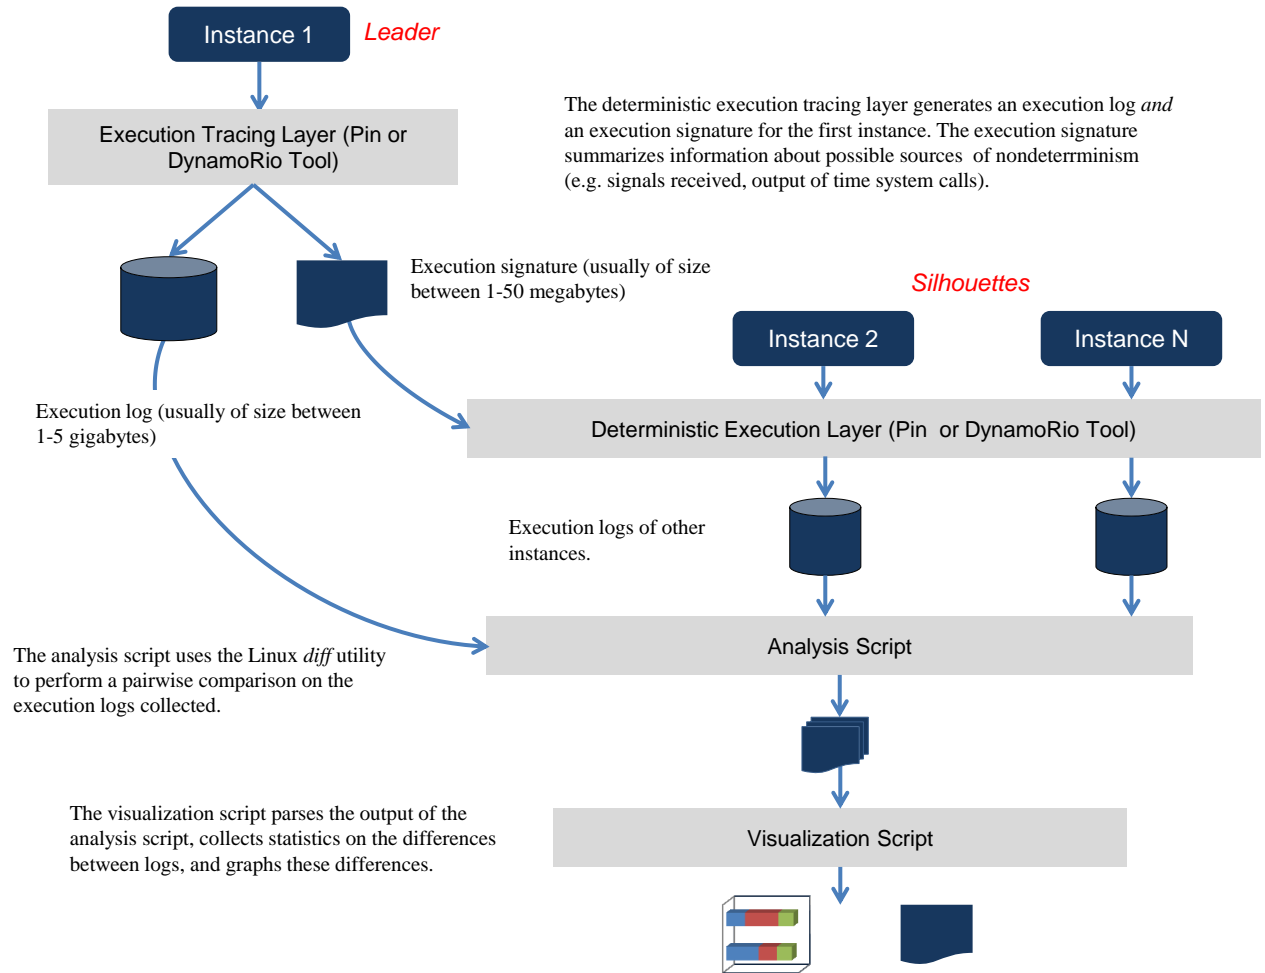


Figure 1-5: Simulation of *Silhouette Execution* in a bootstorm scenario. We use dynamic instrumentation to generate a trace signature file for the leader. While we do not bypass execution in the silhouettes, try to reduce the number of fork-points and record information about them. Our analysis and visualization scripts allow us to simulate and evaluate the effectiveness of *silhouette* execution.

Canary and Pointer Guard Values

The values of the canary (`gs:0x14`) and the pointer guard (`gs:0x18`) are initialized in user-space, so dynamic instrumentation can be used to force these values to agree across distinct executions of the same program: instructions that initialize them can be modified or replaced; the sources used to compute these values (e.g. `rdtsc`, `‘/dev/urandom’` or `AT_RANDOM`) bytes can be intercepted as well.

For a real implementation of silhouette execution, this means that a leader can simply choose to not treat the instructions that initialize `gs:0x14` or `gs:0x18` as fork-points and simply treat them as normal memory writes instead. The silhouettes will follow the leader’s execution signature and store the same canary or pointer guard as the leader into memory. All subsequent instructions that load the canary or pointer guard values from memory will also be identical and thus will not be fork-points.

Randomization

To overcome execution differences resulting from randomization, we need to address the standard techniques used by programs to seed PRNGs. In our simulations, reads performed by the leader from `‘/dev/urandom’`, the `AT_RANDOM` bytes, or the `rdtsc` instruction are intercepted and recorded in the trace execution file using dynamic instrumentation; for other subsequent instances, we simply replace the return values to match those from the leader.

For silhouette execution in practice, this means that the leader can simply exclude `rdtsc` instructions or reads from `‘dev/urandom’` from fork-points and simply execute past them. Semantically, when silhouettes replay the leader’s execution signature, this simulates the unlikely but possible case that they received the same random values as the leader from external sources.

We need a slightly different approach for `AT_RANDOM` bytes because they are not initialized in user-space. Simply excluding reads of `AT_RANDOM` bytes from fork-points is not sufficient for correctness: when execution diverges permanently, a silhouette may read `AT_RANDOM` bytes again and they will be different from those read earlier (which is impossible). To solve this minor issue, we can make the leader transmit its

AT_RANDOM bytes in its first execution signature; the silhouettes overwrite their own AT_RANDOM bytes with the leader’s values before starting execution.

These strategies eliminate any fork-points or tainted memory locations that result from external sources of randomization in programs. While eliminating such randomization can change execution semantics of a Linux service, we are still simulating a valid (and possible) execution path for each silhouette.

Time

In our simulations, system calls that return some measurement of the current time, the CPU or program execution time, or a time interval (e.g. `time` or `gettimeofday`) can be intercepted in the same manner as randomization: the timestamps logged in the trace signature file can be used to force agreement between different instances.

For silhouette execution in practice, this means that a leader can simply exclude such system calls from fork-points, and continue executing past them. Rather than including these system calls in its execution signature, the leader should include the memory side-effects of these system calls. This way, silhouettes do not perform these system calls; rather, they automatically copy the behavior of the leader when they replay its writes. This simulates the unlikely but semantically valid scenario that the various instances executed various time-related system calls precisely at the same times. Replacing such system calls with their memory side-effects works because they do not mutate any operating system state.

The timestamps returned from `stat` system calls are not as easily handled. *If* we assume that all the input, configuration and data files are identical between various instances of a program, then we can simply exclude `stat` system calls from fork-points and include them in the execution signature file instead. A leader can assume by default that only the various timestamps returned by `stat` will be different and mark them as tainted. In our experiences, these timestamps are typically ignored in an overwhelming majority of cases. Thus, tainting these values by default creates little overhead because these values are seldom read or propagated. At the same time, we avoid the overhead associated with treating `stat` system calls as fork-points.

In the rare cases where the timestamps from `stat` system calls are actually read, they are typically compared to determine “freshness” (i.e. which file is newer). When we assume that all the files accessed by a program have similar initial contents, these comparisons can also be excluded from fork-points because timestamps retain the same ordering across different instances of a program in our experiments. This can be a risky optimization, so to be absolutely sure, the leader could treat these comparisons as fork-points and verify that the comparison results are the same in all instances instead.

To model this in our simulations, we do not include `stat` system calls that contain different time-stamps across instances in our count of fork-points. However, comparisons of timestamps from `stat` system calls are included in our definition of fork-points.

Signal Delivery

In order to overcome the unpredictable timing and order of signals in our simulations, we intercept all signals received by an application and ensure they are delivered at precisely the same instruction counts and in the same order as that indicated in the trace signature file. Unlike record-and-replay systems, we only deliver signals that are *actually* received. Thus, signals that are received earlier than expected are simply delayed or reordered. If, however, a signal is not received at the expected instruction count, our instrumentation tool simply waits or inserts `nops` until the desired signal is received. If a signal simply refuses to appear for a long time, execution must diverge. In our experiments, this final case does not occur as long as other sources of nondeterminism are controlled.

For silhouette execution in practice, this means that a leader can exclude received signals from fork-points (unlike before), and simply include them in the execution signature sent to silhouettes instead. When silhouettes bypass execution using the execution signature, they delay, reorder and deliver signals at precisely the same times as the leader and thus avoid any control-flow divergences. This technique, *signal alignment*, requires each silhouette to temporarily resume execution from the

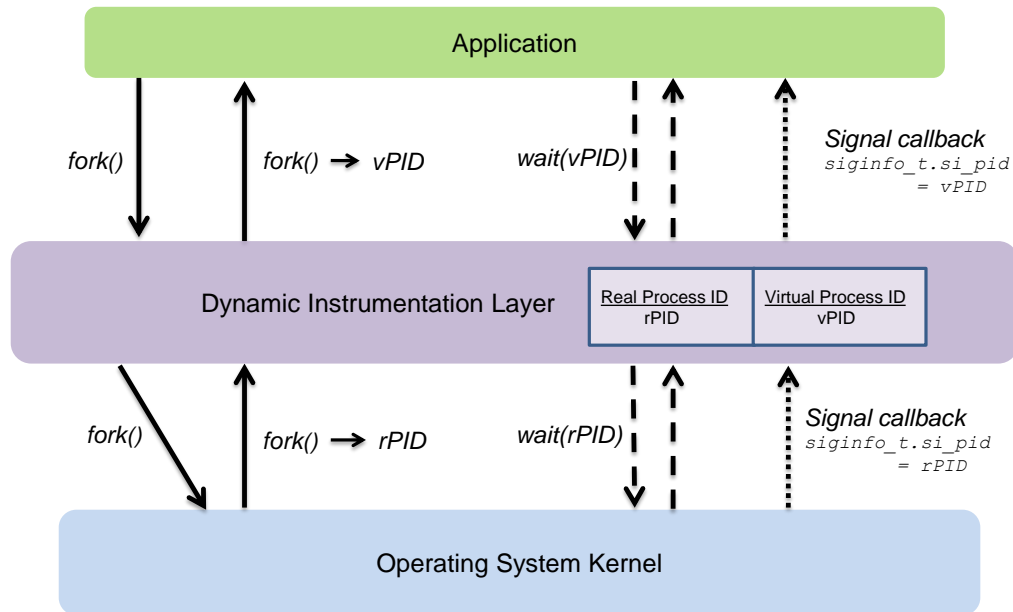


Figure 1-6: All system calls and communications between the Linux user and kernel space are intercepted; the dynamic instrumentation layer uses a PID translation table, and translates between real and virtual process IDs to ensure correctness.

instruction before which the signal is expected to be received, and then insert `nops` until the signal is received or deliver a previously received (and withheld) signal. This can prevent control flow divergences and subsequent fork-points that arise from variable signal timing and ordering. Of course, if the expected signal is not received in a silhouette within a reasonable time interval, control flow must inevitably diverge.

Process IDs

In our simulations, nondeterminism from process IDs can be controlled by *virtualizing the process ID layer*, as shown by Figure 1-6. Using dynamic instrumentation, we can replace real and unpredictable process IDs from kernel space with virtual and unpredictable process IDs in user space. As outlined in Section ??, all interfaces which use process IDs need to be carefully monitored so that process IDs can be translated back and forth for correctness.

In practice, this is equivalent to using an operating system that assigns process IDs in a deterministic fashion. In a complete implementation of silhouette execution, we could expect the leader VM to store process IDs in its execution signature; silhouettes would assign process IDs in the same order as the leader.

For silhouette execution purely in user-space, we could add some logic to the silhouettes and virtualize their process IDs as described in Figure 1-6. Instead of treating instructions that deal with process IDs as fork-points (e.g. `fork` or `clone`), the leader would execute past them, and simply flag them before including them in its execution signature. When silhouettes would arrive at these instructions, they would simply translate between real and virtual process IDs. This clearly adds complexity and translation overhead whenever an instance interacts with an operating system interface using virtual process IDs. However, virtualizing the process ID layer prevents a considerable number of user-space instructions that propagate differences in process IDs (e.g. in `libc` or `pthread`s) from being labeled at fork-points due to taint propagation. Note that, for correctness, silhouettes must pretend to use the same virtual process IDs that they have already used before even after they completely diverge execution.

File I/O

Differences in input file contents across executions would inevitably cause execution to diverge, but reducing fork-points arising from time, randomization or process ID system calls is typically sufficient to ensure that file contents evolve identically in Linux services. Some files that may differ between two instances on start up (e.g. cache files or logs) can simply be deleted or replaced without sacrificing correctness. Also, as mentioned already, `stat` timestamps are frequently not read, so they can be excluded from fork-points until they are actually compared with other timestamps.

For file contents that are inherently and minimally different (e.g. they contain different timestamps), taint propagation otherwise suffices to ensure that they do not cause execution to diverge permanently.

Network I/O

The content of network configuration files does not change over a booting period in our experiments, so we can mark them as immutable during the boot process and arrange for their `stat` timestamps to return identical values across different instances. In practice, this means that a leader can exclude the `stat` calls for network files or their comparisons from its set of fork-points without sacrificing correctness.

In the same vein, the leader can exclude IP address resolution instructions from its definition of fork-points, and execute past them. This disables DNS-based dynamic load balancing during the boot process and forces silhouettes to use the same resolved IP addresses as the leader. DNS-based load balancing typically uses low TTLs, so once the boot process is finished and execution diverges permanently, the IP addresses will be refreshed quickly.

If bytes read from sockets differ across different executions, we need to understand the *context* to determine whether the differences are serious (e.g. due to different requests) or synthetic (e.g. due to timestamps). Simply tainting the bytes that are different between the instances is sufficient to tolerate minor differences from external data. In some specific cases, we can avoid marking socket `reads` as fork-points: for `Netlink` sockets, typically the only differences in bytes across different instances arise from interface statistics (`RTM_NEWLINK`), sequence numbers or end-point IDs. To handle variability in interface statistics (which are rarely read, if at all), we can simply overwrite them with identical fixed values across instances; sequence numbers and end-points IDs are generated by `time` system calls and process IDs respectively, which we already handle. Thus, a leader can execute a socket `read` on `Netlink` sockets and simply add memory operations that replace any statistics read by the silhouette with what the leader read. This simulates the unlikely but semantically valid execution scenario that all instances have the same network statistics (e.g. packets dropped/transmitted) when the system call is made. Using this approach, while we introduce instrumentation overhead in the leader for handling `Netlink` sockets, we reduce the number of expensive fork-points requiring byte-by-byte comparisons related to `Netlink` socket `reads` across all instances.

To overcome fork-points related to ephemeral port assignments, the leader can monitor the `bind` or `connect` system calls and changes their arguments to explicitly request ports in the ephemeral range rather than let the kernel assign them. These system calls, then, are no longer fork-points; the silhouettes derive and use the same arguments as the leader from the execution signature files sent to them, so there are no execution differences. This allows us to avoid virtualizing ephemeral ports in a similar fashion to how we virtualize process IDs.

Scalable I/O Schemes

To handle nondeterminism caused by unpredictable ordering of I/O events, we use techniques similar to those used for reordering signals in our simulations, as described by Figure 1-7. We name our approach *I/O stalling*.

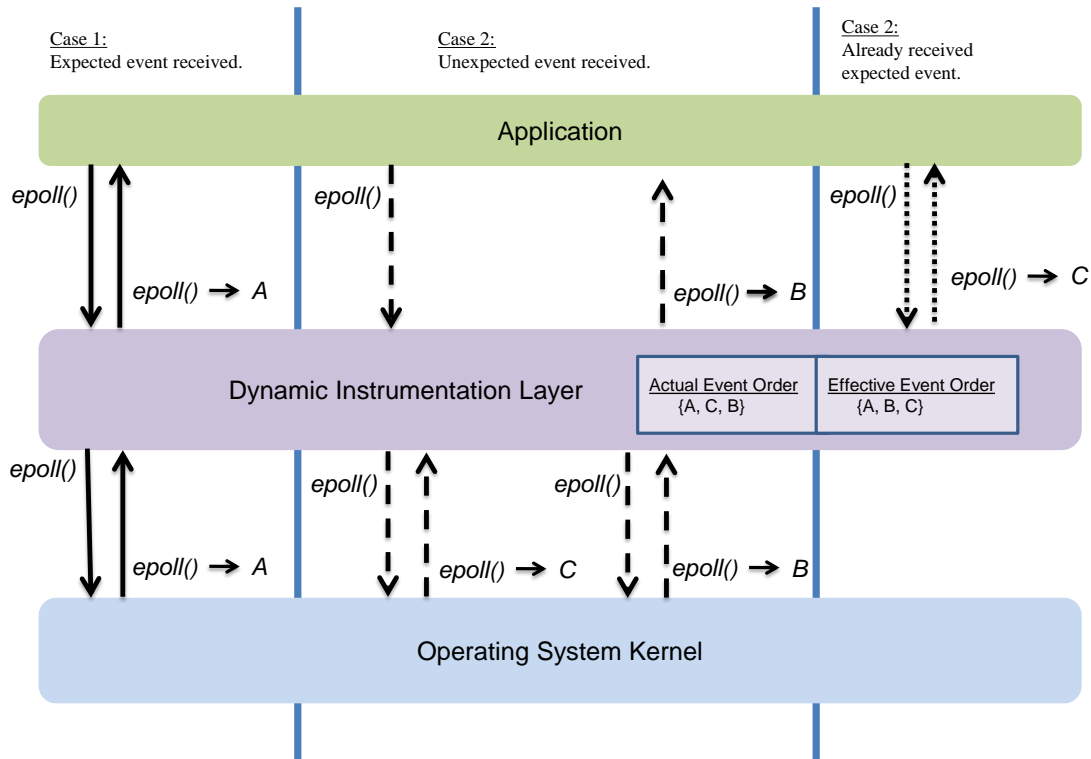


Figure 1-7: We intercept all `epoll` system calls, and use the execution signature file to achieve determinism. We do not “replay” I/O events because only events that actually do occur are delivered to the application instance. This diagram assumes `epoll` returns one event per call for the sake of illustration.

Assuming that `epoll` returns just one event, figure 1-7 illustrates three possible cases that could occur:

- The event returned by a call to `epoll` (A) is the one expected in the execution signature file (A). The instrumentation layer does not modify the system call.
- The desired event (B) has not been received yet, and `epoll` returns an unexpected event (C). The instrumentation layer stores the out-of-order event, and repeatedly calls `epoll` until the the expected event is received.
- A call to `epoll` is initiated, and the event desired (C) has already been received. The instrumentation layer does not make a system call and simulates a return from `epoll` with the expected event instead.

Even if I/O events are reordered, it is possible that different amounts of data are available for ready file descriptors across executions. We can mask this effect in the same way we handle signals: if more bytes are available (e.g. through `read`) than expected in the execution signature file, we modify return values and `read` buffers to delay the reading of these bytes until the next `read`. In some corner cases, we may have to “fake” readiness in a call to `epoll`: if all bytes to be read from a file descriptor have been read by the dynamic instrumentation layer (out of which a few have not yet been delivered to the application), there will be no more readiness events even though the application expects them. If less-than-expected bytes are available, we simply wait till they are available by waiting for another readiness update on the same file descriptor inside dynamic instrumentation layer. In our experiments, this approach has been sufficient for overcoming nondeterminism from event-based I/O schemes. For asynchronous I/O schemes (e.g. `aio_read`), strategies similar to those used for reordering and precisely-timing signals would be necessary to hide variable I/O latency and ordering.

For silhouette execution in practice, we would implement this technique – *I/O alignment* – by excluding `epoll` system calls from fork-points. An `epoll` call and its expected results would be stored in the execution signature; a silhouette would

execute `epoll` and not let differences in the observed results let execution branch immediately. Instead, the silhouette would repeatedly call `epoll` until it gets the desired event set, taking care to delay and reorder I/O events as necessary to align with the leader. The same approach can be used for determining how many bytes are available from a `read`. While this adds overhead to the instructions executed by a silhouette, it eliminates control-flow divergences and subsequent fork-points created from inherent I/O nondeterminism.

Concurrency

Execution variability resulting from multi-threading has been extensively documented; there is a significant body of work that attempts to overcome such nondeterminism by using deterministic logical clocks or record-and-replay approaches. For our experiments, we did not attempt to enforce a total order on the instructions executed in multi-threaded programs and just measured fork-points inside the main process for each Linux service. To reduce fork-points from multi-threading, we could incorporate deterministic logical clocks into our design.

As mentioned before, a nondeterministic system scheduler can cause variable timing of signals or I/O events, which we handle using signal and I/O alignment strategies. Work on deterministic operating systems can be extended to overcome this issue in a more systematic manner.

Procfs: The ‘`/proc/directory`’

A leader can exclude reads from *procfs* for statistics from fork-points; instead, the leader can simply replace them with memory operations that simulate identical results in silhouettes. A leader must be careful to flag reads from *procfs* use process IDs, so that silhouettes can translate between real and virtual process IDs accordingly.

1.5 Evaluation of Improved Silhouette Execution

Bibliography

- [1] T. Bergan, N. Hunt, L. Ceze, and S.D. Gribble. Deterministic process groups in dos. *9th OSDI*, 2010.
- [2] D.L. Bruening. *Efficient, transparent, and comprehensive runtime code manipulation*. PhD thesis, Citeseer, 2004.
- [3] A.T. Clements, I. Ahmad, M. Vilayannur, and J. Li. Decentralized deduplication in san cluster file systems. In *Proceedings of the 2009 conference on USENIX Annual technical conference*, pages 8–8. USENIX Association, 2009.
- [4] J.G. Hansen and E. Jul. Lithium: virtual machine storage for the cloud. In *Proceedings of the 1st ACM symposium on Cloud computing*, pages 15–26. ACM, 2010.
- [5] Solving Boot Storms With High Performance NAS. http://www.storage-switzerland.com/Articles/Entries/2011/1/3_Solving_Boot_Storms_With_High_Performance_NAS.html, 2011. [Accessed 1-August-2011].
- [6] X.F. Liao, H. Li, H. Jin, H.X. Hou, Y. Jiang, and H.K. Liu. Vmstore: Distributed storage system for multiple virtual machines. *SCIENCE CHINA Information Sciences*, 54(6):1104–1118, 2011.
- [7] C.K. Luk, R. Cohn, R. Muth, H. Patil, A. Klauser, G. Lowney, S. Wallace, V.J. Reddi, and K. Hazelwood. Pin: building customized program analysis tools with dynamic instrumentation. In *ACM SIGPLAN Notices*, volume 40, pages 190–200. ACM, 2005.
- [8] Bootchart: Boot Process Performance Visualization. <http://www.bootchart.org>, 2011. [Accessed 29-July-2011].
- [9] S. Meng, L. Liu, and V. Soundararajan. Tide: achieving self-scaling in virtualized datacenter management middleware. In *Proceedings of the 11th International Middleware Conference Industrial track*, pages 17–22. ACM, 2010.
- [10] VMware Bootstorm on NetApp. <http://ctistrategy.com/2009/11/01/vmware-boot-storm-netapp/>, 2009. [Accessed 29-July-2011].
- [11] M. Olszewski, J. Ansel, and S. Amarasinghe. Scaling deterministic multithreading. *2nd WoDet*, 2011.

- [12] H. Patil, C. Pereira, M. Stallcup, G. Lueck, and J. Cownie. Pinplay: a framework for deterministic replay and reproducible analysis of parallel programs. In *Proceedings of the 8th annual IEEE/ACM international symposium on Code generation and optimization*, pages 2–11. ACM, 2010.
- [13] A. Qureshi, R. Weber, H. Balakrishnan, J. Guttag, and B. Maggs. Cutting the electric bill for internet-scale systems. In *Proceedings of the ACM SIGCOMM 2009 conference on Data communication*, pages 123–134. ACM, 2009.
- [14] Vijayaraghavan Soundararajan and Jennifer M. Anderson. The impact of management operations on the virtualized datacenter. *SIGARCH Comput. Archit. News*, 38:326–337, June 2010.
- [15] M.W. Stephenson, R. Rangan, E. Yashchin, and E. Van Hensbergen. Statistically regulating program behavior via mainstream computing. In *Proceedings of the 8th annual IEEE/ACM international symposium on Code generation and optimization*, pages 238–247. ACM, 2010.
- [16] S.B. Vaghani. Virtual machine file system. *ACM SIGOPS Operating Systems Review*, 44(4):57–70, 2010.
- [17] VMware Virtual Desktop Infrastructure. http://www.vmware.com/pdf/virtual_desktop_infrastructure_wp.pdf, 2011. [Accessed 29-July-2011].
- [18] C.A. Waldspurger. Memory resource management in vmware esx server. *ACM SIGOPS Operating Systems Review*, 36(SI):181–194, 2002.